

```
import cv2 # it will allow us to load our images into the script
import numpy as np # used for reformatting our own images
import tensorflow as tf # main library used to load data sets, build neural networks, train them, etc.
import matplotlib.pyplot as plt # used for visualization
```


```
# Load MNIST data
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
```

```
print('The shape of the training inputs:', x_train.shape)
print('The shape of the training labels:', y_train.shape)
print('The shape of the testing inputs:', x_test.shape)
print('The shape of the testing labels:', y_test.shape)
```

```
↗ The shape of the training inputs: (60000, 28, 28)
The shape of the training labels: (60000,)
The shape of the testing inputs: (10000, 28, 28)
The shape of the testing labels: (10000,)
```

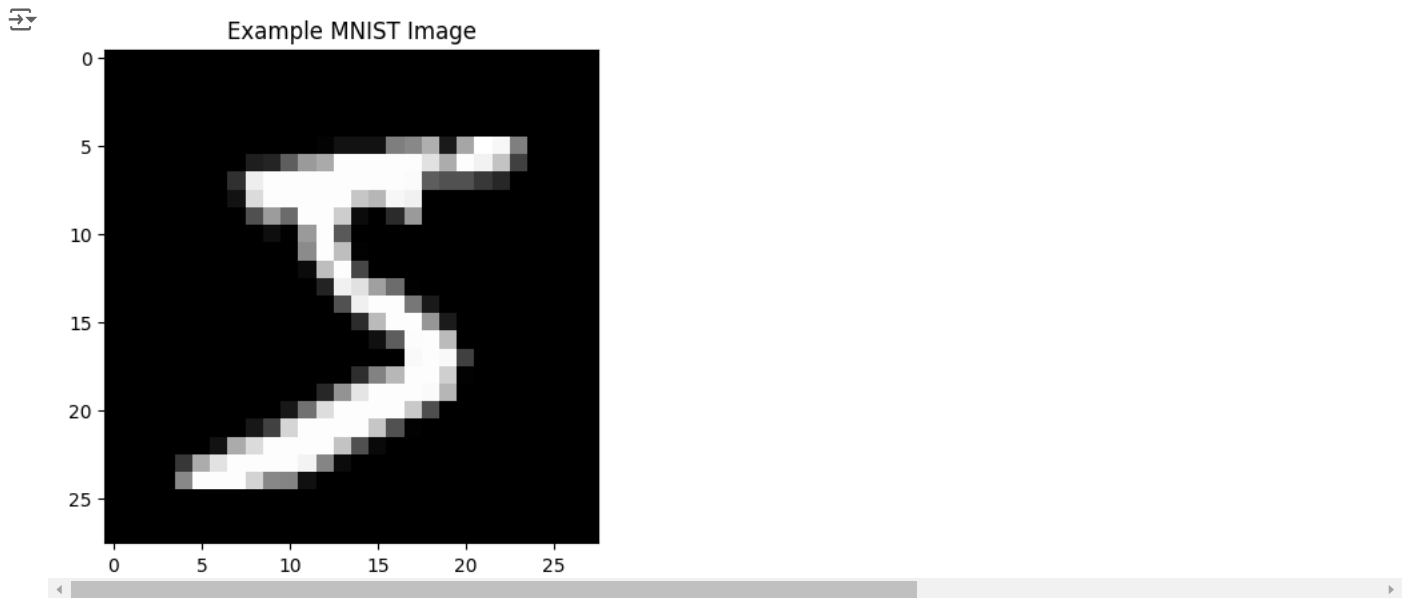
```
# plotting the first 9 images in the train set of MNIST
```

```
fig, axs = plt.subplots(10, 10)
cnt = 0
for i in range(10):
    for j in range(10):
        axs[i, j].imshow(x_train[cnt])
        cnt += 1
```



```
import matplotlib.pyplot as plt
```


```
# Display an example image
plt.imshow(x_train[0].reshape(28,28), cmap='gray')
plt.title('Example MNIST Image')
plt.show()
```



```
import matplotlib.pyplot as plt

# Preprocess the data
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# Display an example image
plt.imshow(x_train[2].reshape(28, 28), cmap='gray') # Reshape the image to 28x28
plt.title('Example Image')
plt.show()
```



```
x_train = tf.keras.utils.normalize(x_train, axis=1)
x_test = tf.keras.utils.normalize(x_test, axis=1)

model = tf.keras.models.Sequential()

model.add(tf.keras.layers.Flatten(input_shape=(28,28)))

/usr/local/lib/python3.10/dist-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input_shape` to `input_shape` in the `__init__` method of `Flatten` layer.
super().__init__(**kwargs)

#Hidden Layer
model.add(tf.keras.layers.Dense(units=128, activation=tf.nn.relu))

model.add(tf.keras.layers.Dense(units=10, activation=tf.nn.softmax)) # output layer

model.summary()
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
flatten_4 (Flatten)	(None, 784)	0
dense_10 (Dense)	(None, 128)	100,480
dense_11 (Dense)	(None, 128)	16,512
dense_12 (Dense)	(None, 10)	1,290

Total params: 118,282 (462.04 KB)

Trainable params: 118,282 (462.04 KB)

Non-trainable params: 0 (0.00 KB)

```
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
# Define a single layer perceptron model
model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Dense(10, activation='softmax', input_shape=(28*28,)))
```

```
# Define a single layer perceptron model
model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Dense(10, activation='softmax', input_shape=(28*28,)))
```

```
# Compile the model - It is important to compile the model after it is defined.
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
# Reshape x_train to flatten the images
x_train_flattened = x_train.reshape(-1, 28*28) # -1 infers the number of samples
```

```
# Train the model using the flattened data
history=model.fit(x_train_flattened,
                  y_train,
                  epochs=10,
                  batch_size=100,
                  validation_split=0.2)
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input\_shape`/`input\_dim` arg to `Dense` layer. It will be ignored. If you want to make sure the layer has the expected input, use `input\_shape` argument in the `model.fit` method.

```
Epoch 1/10
480/480 ————— 4s 5ms/step - accuracy: 0.7142 - loss: 1.3514 - val_accuracy: 0.8798 - val_loss: 0.5113
Epoch 2/10
480/480 ————— 4s 4ms/step - accuracy: 0.8724 - loss: 0.5048 - val_accuracy: 0.8988 - val_loss: 0.3919
Epoch 3/10
480/480 ————— 2s 2ms/step - accuracy: 0.8922 - loss: 0.3981 - val_accuracy: 0.9062 - val_loss: 0.3500
Epoch 4/10
480/480 ————— 2s 4ms/step - accuracy: 0.9018 - loss: 0.3612 - val_accuracy: 0.9120 - val_loss: 0.3274
Epoch 5/10
480/480 ————— 2s 2ms/step - accuracy: 0.9043 - loss: 0.3384 - val_accuracy: 0.9142 - val_loss: 0.3143
Epoch 6/10
480/480 ————— 1s 2ms/step - accuracy: 0.9086 - loss: 0.3250 - val_accuracy: 0.9153 - val_loss: 0.3053
Epoch 7/10
480/480 ————— 1s 2ms/step - accuracy: 0.9141 - loss: 0.3057 - val_accuracy: 0.9172 - val_loss: 0.2977
Epoch 8/10
480/480 ————— 1s 2ms/step - accuracy: 0.9150 - loss: 0.3020 - val_accuracy: 0.9183 - val_loss: 0.2934
Epoch 9/10
480/480 ————— 1s 2ms/step - accuracy: 0.9170 - loss: 0.2947 - val_accuracy: 0.9187 - val_loss: 0.2890
Epoch 10/10
480/480 ————— 1s 2ms/step - accuracy: 0.9189 - loss: 0.2896 - val_accuracy: 0.9183 - val_loss: 0.2871
```

```
# Evaluate the model
x_test_flattened = x_test.reshape(-1, 28*28) # Flatten the x_test data
test_loss, test_acc = model.evaluate(x_test_flattened, y_test)
print(f'Test accuracy: {test_acc}')
```

```
313/313 ————— 1s 2ms/step - accuracy: 0.9077 - loss: 0.3242
Test accuracy: 0.920199990272522
```

```
from sklearn.metrics import confusion_matrix, classification_report
import numpy as np
```

```
# Generate predictions using the model
y_pred = model.predict(x_test_flattened)
y_pred_classes = np.argmax(y_pred, axis=1) # Convert predictions to class labels
```

```
# Generate the confusion matrix
cm = confusion_matrix(y_test, y_pred_classes)
print("Confusion Matrix:")
print(cm)
```

```
# Generate the classification report
cr = classification_report(y_test, y_pred_classes)
print("\nClassification Report:")
print(cr)
```

313/313 1s 2ms/step

Confusion Matrix:

```
[[ 950   0   0   1   0  10  10   3   5   1]
 [   0 1111   4   1   0   3   4   0  12   0]
 [   9   9 909  21  12   2  13  11  42   4]
 [   3   0  14 920   1  25   5   8  23  11]
 [   1   1   6   1 911   0  14   4   8  36]
 [   9   1   6  37   5 754  20   7  45   8]
 [  12   3   7   1   8  11 910   1   5   0]
 [   2  10  26   7   6   0   0 938   2  37]
 [   6   6   5  22  11  26  10  10 866  12]
 [   9   6   2  12  17   5   0  15  10 933]]
```


Classification Report:

	precision	recall	f1-score	support
0	0.95	0.97	0.96	980
1	0.97	0.98	0.97	1135
2	0.93	0.88	0.90	1032
3	0.90	0.91	0.91	1010
4	0.94	0.93	0.93	982
5	0.90	0.85	0.87	892
6	0.92	0.95	0.94	958
7	0.94	0.91	0.93	1028
8	0.85	0.89	0.87	974
9	0.90	0.92	0.91	1009
accuracy			0.92	10000
macro avg	0.92	0.92	0.92	10000
weighted avg	0.92	0.92	0.92	10000

```
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np # Added import for numpy



# Generate the confusion matrix
y_pred_classes = np.argmax(y_pred, axis=1) # Convert predictions to class labels
cm = confusion_matrix(y_test, y_pred_classes) # Use class labels for y_pred

# Plot the confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix")
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.show()
```



```
import matplotlib.pyplot as plt
# Plot training & validation accuracy values
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Plot training & validation loss values
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



```
import numpy as np
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize
import matplotlib.pyplot as plt

# Get predicted probabilities for each class
# Reshape x_test to be 2-dimensional
x_test_flattened = x_test.reshape(-1, 28*28)
y_pred_proba = model.predict(x_test_flattened)

# Binarize the labels
y_test_bin = label_binarize(y_test, classes=np.arange(10))

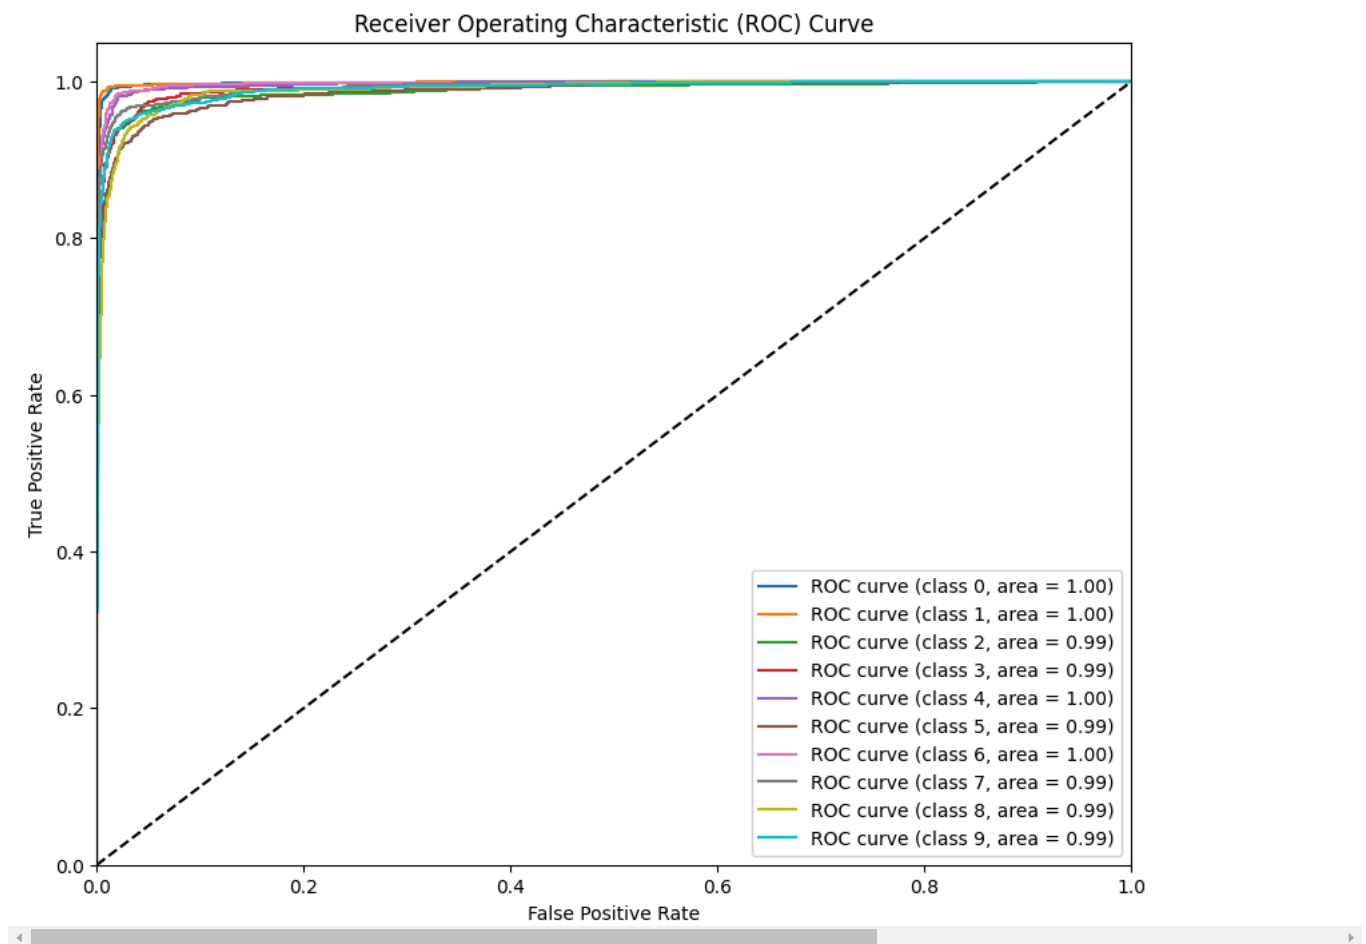
# Compute ROC curve and AUC for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(10):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_pred_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot ROC curve for each class
plt.figure(figsize=(10, 8))
for i in range(10):
    plt.plot(fpr[i], tpr[i], label='ROC curve (class %d, area = %0.2f)' % (i, roc_auc[i]))

plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
```

```
plt.show()
```

313/313 27s 87ms/step



```
import matplotlib.pyplot as plt
import numpy as np
# Select a random image from the test dataset
image_index = np.random.randint(0, len(x_test))
test_image = x_test[image_index]

# Reshape the image to match the model's input shape
test_image = test_image.reshape(1, 28*28) # Changed to flatten the input image

# Make a prediction
prediction = model.predict(test_image)

# Get the predicted class
predicted_class = np.argmax(prediction)

# Display the image and the prediction
plt.imshow(test_image.reshape(28, 28), cmap='gray')
plt.title(f'Predicted Class: {predicted_class}')
plt.show()
```

 1/1 ————— 0s 38ms/step

Predicted Class: 1

0

```
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import confusion_matrix, classification_report

#Import the models module from keras
from tensorflow import keras
from keras import models
from keras import layers

# Define the MLP model
model = models.Sequential()
model.add(layers.Flatten(input_shape=(28, 28))) # Flatten the input images
model.add(layers.Dense(128, activation='relu')) # Hidden layer 1
model.add(layers.Dense(64, activation='relu')) # Hidden layer 2
model.add(layers.Dense(10, activation='softmax')) # Output layer
```